

# Dynamic Memory Allocation

(summary)

# Dynamic memory allocation

- Ask for memory, as we need it.
- Biggest limitation: We must free it ourselves once done with it.
- Uses:
  - Dynamic array allocation - when we don't know input size
  - Large array / data structure - start small, increase size as we need it
  - Other data structures - Linked lists, etc.

# Dynamic memory allocation example

Build a program that reads in a list of integers, and outputs its sum.

The program should handle lists of **any** (reasonable) **length**.

# Dynamic memory allocation example

Build a program that reads in a list of integers, and outputs its sum.

The program should handle lists of **any** (reasonable) **length**.

How to do this?

1. Use a variable to hold the *capacity* of the array :: *size*
2. Start with 'small' capacity - say 10.
3. Use **malloc** to allocate a block of memory of that size to hold the array.
4. Read in each integer once at a time.
5. If we read in *size* integers, **double** *size* and **realloc**.
6. Finish reading the integers
7. Do our required task -- calculate sum, output them.
8. Release the block of memory allocated with **free**.

# Dynamic memory allocation - <stdlib.h>

`void *malloc(size_t size)`

Allocates a contiguous block of memory of SIZE bytes, returning a void Pointer to the first byte. Elements are left *uninitialised*.

`void free(void *ptr)`

Frees the memory block pointed to by ptr, IF that pointer once returned by malloc / calloc / realloc. Pointer address is *NOT* changed.

`void *calloc(size_t nmemb, size_t size)`

Same as malloc, but (1) checks if nmemb \* size will integer overflow, and (2) initialises all elements to 0. (Slow)

`void *realloc(void *ptr, size_t size)`

Changes size of memory block pointed to by ptr to one having size SIZE, returning a pointer to that memory block. Some caveats.

# When to use Dynamic Memory Allocation

- No upper bound on size of input
- There is an upper bound on size of input, but it's large -> waste memory
- Want to use a lot of memory
- Want to have memory / data structures persist in your program
- Want memory on the heap

# Dynamic memory allocation issues.

Gotta watch out when doing dynamic memory allocation.

Issues to consider:

- Failed to allocate memory
- Double Free
- Memory Corruption

'High' address 0xFFFF FFFF

*(Command-line arguments  
Environment variables)*

```
int x[10];
```

**stack** local variables

```
int *A = malloc(10*sizeof(int));
```

**heap** dynamic memory allocation  
(malloc)

```
static int p = 5;  
int global = 4;
```

**data** static, global variables  
function pointers

**code** program code (binary)

'Low' address 0x0000 0000



<b>Memory allocation:</b>	<b>Automatic</b>	<b>Static</b>	<b>Dynamic</b>
<i>Managed by...</i>	the compiler	the compiler	the programmer
<i>Alloc'd...</i>	during variable initialisation	at compile time	using malloc / calloc / realloc
<i>Free'd...</i>	automatically, when exiting scope of variable	automatically, at end of program	manually, when free is called <i>!! memory leaks !!</i>
<i>Memory assigned to...</i>	<b>stack</b> , within <b>stack frame</b> (at runtime)	<b>data</b> (at compile time)	<b>heap</b> (at runtime)
<i>Examples / Uses</i>	local variables, function arguments, array size known	global, static variables	Dynamically sized arrays, structs, array size unknown, Large data structures, custom data structures
<i>Issues</i>	Limited control of lifetime; Can't change size after init; Size calculated at compile time; Size limits	Variable is allocated permanently, wastes memory; Same limitations as automatic	Memory leak (no free); Double free; Additional memory required to store pointer to memory;

# Static & Automatic allocation

# Scenario

Build a program that reads in a list of integers, and outputs its sum.

The program should handle lists of **any** (reasonable) **length**.

# Memory Allocation

- **Automatic**
  - By the compiler, at runtime, stored in **stack**
- **Static**
  - By the compiler, at compile time, stored in **data**
- **Dynamic**
  - By the programmer, at runtime, stored in **heap**

'High' address 0xFFFF FFFF

*Recall memory layout*

*(Command-line arguments  
Environment variables)*


**stack** local variables

**heap** dynamic memory allocation  
(malloc)

**data** static, global variables  
function pointers

**code** program code (binary)

'Low' address 0x0000 0000



# Static & Automatic Memory Allocation

(1)

What's in memory, and what's in scope at each point?

```
int g = 4;
void func(int arg, char *str) {
    static int x = 5; (2)
    double pair[2] = {1.4, 4.3};
    for (int i = 0; i < g; i++) {
        printf("%d ", i * x++); (3)
    }
} (4)
```

(5)

# Static & Automatic Memory Allocation

```
int g = 4;
void func(int arg, char *str) {
    static int x = 5;
    double pair[2] = {1.4, 4.3};
    for (int i = 0; i < g; i++) {
        printf("%d ", i * x++);
    }
}
```

## Automatic memory allocation

- automatically allocated by the compiler at runtime, once we enter scope of the variable
- freed once we exit the scope of the variable.
- stored within **stack**

'High' address 0xFFFF FFFF

*Where do those  
variables go?*

*(Command-line arguments  
Environment variables)*


**stack** local variables

**heap** dynamic memory allocation  
(malloc)

**data** static, global variables  
function pointers

**code** program code (binary)

'Low' address 0x0000 0000





# Static & Automatic Memory Allocation

```
int g = 4;
void func(int arg, char *str) {
    static int x = 5;
    double pair[2] = {1.4, 4.3};
    for (int i = 0; i < g; i++) {
        printf("%d ", i * x++);
    }
}
```

## Static memory allocation

- allocated by the compiler at compile time, **once**.
- freed at program exit.
- stored within **data**

# Problems with **automatic** allocation

- Size calculated at compile time
  - no dynamically sized arrays
  - eg: Can't create array based on input size
- Allocated size can't be changed after initialisation
  - Need to allocate upper bound of requirement -- wastes memory. eg: storing between 10-10000 items, must create array of size 10000 to ensure there's enough space
- Limited stack memory
  - Limited size to memory blocks
  - "Array too large"

# Problems with **static** allocation

- Once a variable is allocated, it is permanently allocated
  - Free'd at program exit
  - Waste memory
- Created at compile time -- can only have one of each item
- Same problems as automatic allocation

# Scenario

Build a program that reads in a list of integers, and outputs its sum.

The program should handle lists of **any** (reasonable) **length**.

Can we do this with automatic / static allocation?

# stdlib.h memory functions

# Dynamic memory allocation - <stdlib.h>

`void *malloc(size_t size)`

Allocates a contiguous block of memory of SIZE bytes, returning a void Pointer to the first byte. Elements are left *uninitialised*.

`void free(void *ptr)`

Frees the memory block pointed to by ptr, IF that pointer once returned by malloc / calloc / realloc. Pointer address is *NOT* changed.

`void *calloc(size_t nmemb, size_t size)`

Same as malloc, but (1) checks if nmemb \* size will integer overflow, and (2) initialises all elements to 0. (Slow)

`void *realloc(void *ptr, size_t size)`

Changes size of memory block pointed to by ptr to one having size SIZE, returning a pointer to that memory block. Some caveats.

# void \*malloc(size\_t size)

Allocates a contiguous block of memory of SIZE bytes within the heap, returning a void pointer to the first byte. Elements are left *uninitialised*.

```
int *arr = (int *) malloc(5 * sizeof(int));
```

^ Allocates a block of size 5 ints within the heap. Stores an address in the stack to the first element of that block in arr. Effectively the same as:

```
int arr[5];
```

Difference? Malloc: we need a pointer to access the memory block - an additional 8B of memory.

```
arr[0] = 5;      works for both.
```

# void \*malloc(size\_t size)

*Equivalencies:*

```
int *arr = (int *) malloc(5 * sizeof(int));
```

```
int *arr = malloc(5 * sizeof(int));
```

Not necessary to cast, but may make your code clearer.

```
int *arr;
```

```
... some code ...
```

```
arr = malloc(5 * sizeof(int));
```

// what is the type of arr?

```
// OR
```

```
arr = (int *) malloc(5 * sizeof(int)); // ah, it's an int *
```

```
arr = malloc(5 * sizeof(*arr));
```

// sizeof takes variables



```
void *malloc(size_t size)
```

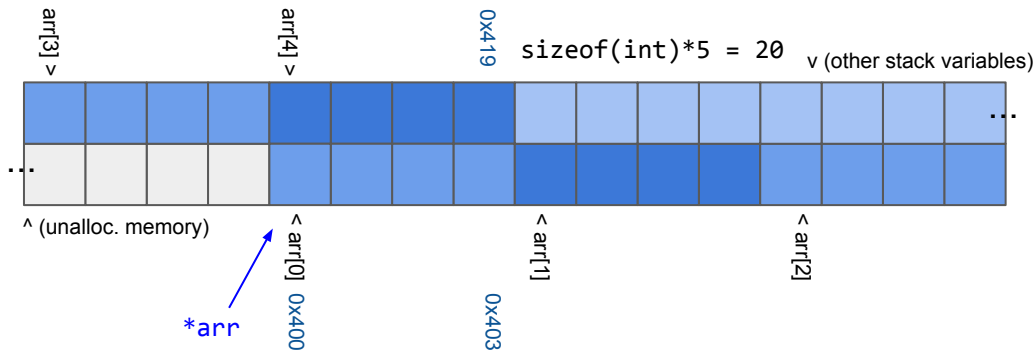
```
int arr[5];
```

 Stack

```

?? ...
[0x400 / stack] arr[0] = uninit
                  arr[1] = uninit
                  arr[2] = uninit
                  arr[3] = uninit
                  arr[4] = uninit
                  ?? ...

```




# void \*malloc(size\_t size)


```
int *arr = (int *) malloc(5 * sizeof(int));
```

Stack Heap

[0x400 / stack] arr = 0x100



?? ...  
[0x100 / heap] arr[0] = uninit  
arr[1] = uninit  
arr[2] = uninit  
arr[3] = uninit  
arr[4] = uninit  
?? ...



(graphics budget exhausted)

# void \*malloc(size\_t size)

We wish to create an int array of size 5 on the heap.  
Which is the correct malloc call?

```
// (a)  
int *arr = (int *) malloc(5 * sizeof(*arr));
```

```
// (b)  
int *arr = (int *) malloc(5 * sizeof(arr));
```

# Dynamic memory allocation - <stdlib.h>

**void \*malloc(size\_t size)**

Allocates a contiguous block of memory of SIZE bytes, returning a void Pointer to the first byte. Elements are left *uninitialised*.

**void free(void \*ptr)**

Frees the memory block pointed to by ptr, IF that pointer once returned by malloc / calloc / realloc. Pointer address is *NOT* changed.

**void \*calloc(size\_t nmemb, size\_t size)**

Same as malloc, but (1) checks if nmemb \* size will integer overflow, and (2) initialises all elements to 0. (Slow)

**void \*realloc(void \*ptr, size\_t size)**

Changes size of memory block pointed to by ptr to one having size SIZE, returning a pointer to that memory block. Some caveats.

# void free(void \*ptr)

Frees the memory block pointed to by ptr, IF that pointer once returned by malloc / calloc / realloc. Pointer address is *NOT* changed.

When we call malloc, we **MUST** free it later.

```
int *arr = malloc(n * sizeof(int));
```

```
// do stuff with arr
```

```
free(arr);    // arr is free'd
```

Some things to be careful about ....

# `void *calloc(size_t nmemb, size_t size)`

Same as malloc, but:

- (1) checks if `nmemb * size` will integer overflow, and
- (2) initialises all elements to 0.

Because of (2) - this can be slow!

```
int *x = calloc(10000, sizeof(int));
```

Returns a pointer to an allocated block of memory to store 10000 ints.

All ints are set to 0. => 10000 operations! Necessary??

```
void *calloc( size_t nmemb, size_t size) { // (roughly)
    if (nmemb > MAX_INT / size) return; // size will overflow
    void *p = malloc(nmemb * size);
    return memset(p, 0, nmemb*size);
}
```

# memset

```
void *memset(void *ptr, int c, size_t n)
```

Sets the first **n** bytes of memory at **ptr** to value **c**.

```
memset(p, 0, n);
```

Same as:

```
int arr[] = {0, 0, 0, 0, 0, 0};    // for n = 6
```

# `void *realloc(void *ptr, size_t size)`

Changes size of memory block pointed to by ptr to one having size SIZE, returning a pointer to that memory block.

## Behaviour

If `ptr == NULL`: `== void *malloc(size_t size)`

If `size == 0 && ptr == NULL`: `== void free(void *ptr)`

If `size > original size`: the new elements are uninitialised

If `size < original size`: some elements are removed / truncated



Google Chrome's main method leaked

```
int main() {  
    void *v = malloc(600000000000);  
    launchchrome();  
    free(v);  
    return 0;  
}
```

# Dynamic allocation issues

# Dynamic memory allocation issues.

Gotta watch out when doing dynamic memory allocation.

Issues to consider:

- Failed to allocate memory
- Double Free
- Memory Corruption

Demonstration. Slides ahead for your own reading!

# When malloc fails

malloc can fail (rarely) when no memory can be allocated.

- System is out of memory (embedded systems have <100mb ram)
- No contiguous free area of requested size can be located in the heap

This happens rarely - but you **should** still take care of it.

When malloc fails, it returns NULL. NULL = (void \*) 0  
(void pointer with value 0)

```
char *str = malloc(INIT_SIZE * sizeof(char)); // sizeof(char) == 1
if (str == NULL) { // handle memory allocate fail
    fprintf(stderr, "Failed to allocate memory for str!");
    exit(EXIT_FAILURE);
}
// ... do stuff with str
```

# Detect malloc fail - variants

```
char *str = malloc(INIT_SIZE * sizeof(char));

if (str == NULL) {
    fprintf(stderr, "Failed to allocate memory for str!");
    exit(EXIT_FAILURE);
}
// ... do stuff with str
```

# Detect malloc fail - variants

```
char *str = malloc(INIT_SIZE * sizeof(char));

if (!str) {
    fprintf(stderr, "Failed to allocate memory for str!");
    exit(EXIT_FAILURE);
}
// ... do stuff with str
```

# Detect malloc fail - variants

```
char *str = malloc(INIT_SIZE * sizeof(char));
```

```
assert(str);
```

```
// do stuff with str
```

# Detect malloc fail - variants

A useful library function to create:

```
void check_ptr(void *p, char *msg) {  
    if (p == NULL) {  
        fprintf(stderr, "check_ptr: %s", msg);  
        // .. do other stuff ..  
        exit(EXIT_FAILURE);  
    }  
}
```

```
char *str = malloc(INIT_SIZE * sizeof(char));  
check_ptr(str, "str alloc failed");
```

```
// do stuff with str
```



# Double Free

```
void func() {  
    int *arr = malloc(n * sizeof(int));  
  
    // do stuff with arr  
  
    // we're done with arr!  
    free(arr);  
  
    // do some other stuff  
  
    // ah, let's make sure we free'd all our malloc'd variables  
    free(arr); // bad crash.  
}
```

\*\*\* Error in `./program': double free or corruption (fasttop): 0x085f6008 \*\*\*

# Double Free

```
void func() {  
    int *arr = malloc(n * sizeof(int));  
  
    // do stuff with arr  
  
    // we're done with arr!  
    free(arr);  
    arr = NULL;  
    // do some other stuff  
  
    // ah, let's make sure we free'd all our malloc'd variables  
    free(arr);    // nothing happens  
}
```

# Memory corruption

```
void func() {  
    int *arr = malloc(n * sizeof(int));  
  
    // do stuff with arr  
  
    // we're done with arr!  
    free(arr);  
  
    arr[4] = 4;    // BAD! that's not malloc'd anymore.  But no error!  
}
```

A **silent** error! Be careful! Very hard to debug!

A solution?

# Memory corruption

```
void func() {  
    int *arr = malloc(n * sizeof(int));  
  
    // do stuff with arr  
  
    // we're done with arr!  
    free(arr);  
    arr = NULL;  
    arr[4] = 4;    // NULL[4] <- Segmentation fault (core dumped)  
}
```

A segmentation fault is still bad, but at least we can debug it. The program will simply crash on that line.

# free issues: a more realistic scenario

```
void func() {  
    int *arr = malloc(n * sizeof(int));  
    int *p = arr;    // a copy of that array pointer  
    // do stuff with arr  
    // we're done with arr! Let's free it  
    free(arr);  
  
    // do some other stuff  
    // do stuff with p -> memory corruption (silent error)  
    // done with p! Let's free it  
    free(p);    // Double free (hard crash)  
}
```

Program would crash on free(p), but if we deleted it, we'd still have memory corruption!

\*\*\* Error in `./program': double free or corruption (fasttop): 0x085f6008 \*\*\*

```
// Q2    :: do this one last!
```

Now: this is an efficient solution. DON'T do an efficient solution!

```
#include <stdlib.h> // malloc, realloc
```

```
char *compute_interleave(char *strs[], int k) {
```

```
    // endstr[i] tells us if we've reached the end
```

```
    // of string strs[i], so shouldn't add \0 or garbage
```

```
    // The end of strs[i] is not necessarily \0s!
```

```
    // can buffer overflow!
```

```
    char *endstr = malloc(k);
```

```
    memset(endstr, 1, k); // or use a for loop
```

```
    int strslength = k;
```

```
    // need to allocate memory properly!
```

```
    int len = 0, size = 2*k;
```

```
    char *retstr = malloc(size); // don't need sizeof -- char
```

```
    // strs[i][j].
```

```
    // strs is an array of strings
```

```
    // strs[0] is the first string
```

```
    // strs[0][0] is the first char of the first string
```

```
    for (int j = 0; strslength; j++) {
```